

**Java course - IAG0040**

# **I/O, Files & Streams**

# File System

- **java.io.File** provides system-independent view of hierarchical pathnames (immutable)

```
-File f = new File("/bin");  
  f = new File(f, "ping");
```

- Can be used to represent files or directories
  - check for existence and permissions
  - query various info (length, attributes)
  - Create, rename or delete both files and directories
  - static fields provide quick access to system-dependent separators: `File.separator`, `File.pathSeparator`

- ' / ' works on all platforms, including Windows

Lecture 8 Slide 2

Java course - IAG0040 Anton Kek

# File System Tips

- How to avoid dealing with separators

- `File parent = new File("someDir");`
  - `File subdir = new File(parent, "bin");`

- Obtain a valid temporary file

- `File tmpFile = File.createTempFile("something", ".tmp");`
  - `tmpFile.deleteOnExit();`

- Enumerate Windows drives

- `File[] roots = File.listRoots();`
  - `File unixRoot = roots[0];`

- Enumerate files in a directory

```
-File[] files = new File("someDir").listFiles();
```

## Lecture 8 Slide 3

Java course - IAG0040 Anton Kek

# Random Access

- `java.io.RandomAccessFile`

- is used when streams are not suitable, e.g. random access of large binary files

- like a large array of bytes in the file system

- both reading and writing is possible (depending on mode)

- Has a file pointer that can be read and moved -

- `getFilePointer()`, `seek()`

- Reading/writing methods are very similar to various

# Input/OutputStreams

– even `DataInput` and `DataOutput` are implemented

Lecture 8 Slide 4

Java course - IAG0040 Anton Kek

## I/O as Streams





Java course - IAG0040 Anton Keks  
Lecture 8 Slide 5

# I/O and Streams

- Java provides I/O facilities in 2 packages
  - `java.io` - traditional synchronous
  - `java.nio` - stream-based I/O
- 'new' (asynchronous)

block-based I/O

- I/O is about reading and writing data
- Reading and writing is possible using

**Streams** - Streams are processed sequentially

- Streams are independent of nature of data
- Java provides

two types (hierarchies) of  
Streams – Byte streams  
(binary)

– Character streams (unicode  
text)

Java course - IAG0040  
Anton Keks

# processing

Reading

## Basic Stream

- open a stream  
(defines the source) •

while more information –

read information

- close the stream

Provided by:

Writing

- `java.io.InputStream`

- `java.io.Reader`

Java course - IAG0040

Anton Kek

- open a stream  
(defines the destination) •

while more information –

write information



- close the stream

- java.io.Writer

Lecture 8  
Slide 7

Provided by:

- java.io.OutputStream

## Basic operations

- Reader and InputStream

- `int read()` - reads a single byte/character (returned as int)

- `int read(byte/char buf[])` - reads as many bytes as possible into the passed array, returns number of bytes read

- `int read(byte/char buf[], int offset, int length)` - the same, but works with the specified portion of an array

- In addition, both support marking locations within a stream, skipping input data, and resetting current position

- **Writer and OutputStream**

- `void write(...)` methods are analogous with reading

- `void flush()` - flushes all buffered data to the output (if any)

**Lecture 8 Slide 8**

Java course - IAG0040 Anton Keks

# Opening and Closing

- Streams are automatically opened on creation - If you have a stream instance - it is ready for reading or writing
- Closing is explicit

- `close()` method closes the stream

- `flush()` is implicit during closing of output streams

- Frees all underlying resources
- Is not the same as object destruction
- Always call it as early as possible
- After `close()` is called, any reads or writes will fail
- Closing several times is safe

Lecture 8 Slide 9

Java course - IAG0040 Anton Kek

# IOException

- I/O classes throw checked *IOExceptions*
- Used by both `java.io` and `java.nio`
- There are many more specific derived

exceptions, like *FileNotFoundException*, *EOFException*, *CharacterCodingException*, etc

- Even the `close()` method can throw an *IOException* (unfortunately)

– Leads to TCFTC (try-catch-finally-try-catch)

Lecture 8 Slide 10

Java course - IAG0040 Anton Keks

## Stream classification

- **Direction:** input and output
- **Data type:** binary and character
- Sink streams or 'endpoints'

- *FileInputStream, ByteArrayInputStream, StringReader, etc*
- Processing streams (wrappers)
  - Base classes: *FilterInputStream, FilterOutputStream, FilterReader, FilterWriter*
  - *SequenceInputStream, ObjectInputStream, BufferedInputStream, LineNumberReader, etc*
- Bridges from binary to characters
  - *InputStreamReader, OutputStreamWriter*

Lecture 8 Slide 11

Java course - IAG0040 Anton Keks

## In-memory I/O

- These streams operate on in-memory data structures, which are passed to them on

creation

- *ByteArrayInputStream, ByteArrayOutputStream*
- *CharArrayReader, CharArrayWriter*
- *StringReader, StringWriter*
- *StringBufferInputStream* (deprecated)
- Useful for mocking streams

Lecture 8 Slide 12

Java course - IAG0040 Anton Keks

## File I/O

- Reading files

- *FileInputStream* - reads binary files
- *FileReader* - reads text files using the default encoding
- *InputStreamReader* can be used for other encodings •

Writing files

- *FileOutputStream* - writes binary files
- *FileWriter* - writes text files using the default encoding
- Task:
  - write a simple 'copy' program (*SimpleCopyProgram* class), implement *net.azib.java.lessons.io.FileCopier*

Lecture 8 Slide 13

Java course - IAG0040 Anton Keks

# Buffering

- These streams wrap other streams to provide buffering capabilities

- *BufferedInputStream, PushbackInputStream*

- *BufferedOutputStream*

- *BufferedReader, PushbackReader*

- *BufferedWriter*

- Task:

- write *BufferedCopyProgram* (implementing *FileCopier*)

- measure performance of both implementations with  
`System.currentTimeMillis()`

**Lecture 8 Slide 14**



# Formatted printing

- Provide convenient printing (e.g. to the console, file, another *Writer*, or *OutputStream*)
- Write-only
  - *PrintWriter* (is preferred)
  - *PrintStream* (*System.out* and *System.err*)
- Often other writable streams are wrapped into these
- They do internal buffering, either a newline char or special method invocations automatically flush the data in case *autoFlush* is enabled
- Warning: *IOExceptions* are never thrown out!
  - `checkError()` may be used for error checking

# Misc features

- Concatenation:

- *SequenceInputStream* - allows to concatenate multiple *InputStreams* together

- Pipes:

- *PipedInputStream*, *PipedOutputStream*, *PipedReader*, *PipedWriter* - allows to connect *InputStream* to an *OutputStream*

- Counting:

- *LineNumberReader* - counts line numbers while reading

- Peeking ahead:

- *PushbackInputStream* and *PushbackReader* - allows to return read data back to stream

- Arbitrary data:

– *DataStream, DataOutputStream* - allows to read/write primitive types easily

## Lecture 8 Slide 16

Java course - IAG0040 Anton Keks

# Serialization

- Java natively supports serialization of data (persistent storage of objects' internal state)
  - Serialization: *ObjectOutputStream*
  - Deserealization: *ObjectInputStream*
- Interfaces
  - *Serializable* - marker but has some methods documented
  - *Externalizable* - defines methods for saving/restoring data manually during the serialization
- It is highly recommended to define `static final long serialVersionUID` field in serializable classes (used for compatibility checks), otherwise it is computed automatically

- fields, declared as **transient** or **static**, are not serialized
- Can be used for **deep** cloning, faster than cloning recursively

Java course - IAG0040 Anton Keks

Lecture 8 Slide 17

# Channels and Buffers

- **java.nio** offers a new way for I/O: Channels and Buffers
  - All data operations are performed on Buffers (data blocks)
    - There are buffers for all primitive types, like *ByteBuffer*, *LongBuffer*, *FloatBuffer*, etc
    - Buffers are allocated using their static methods
    - Buffers are internally arrays, which maintain their capacity, limit, and position, thus simplifying writing code

- Channels are like Streams, but they are bi-directional and read or write data from/into Buffers only. No direct access to data is possible.

Lecture  
8 Slide  
18

# Buffers

- All buffers are mutable
- *ByteBuffer* is most used, provides methods for reading/writing all other primitive types, absolute (index-based) and relative (current position based)
- *ByteBuffer* provides views as other buffers, e.g. *IntBuffer*, *FloatBuffer*
- Direct buffers are backed by OS native storage if possible, non direct buffers are Java arrays

`ByteBuffer.allocate()` - allocates a non-direct buffer

`ByteBuffer.allocateDirect()` - allocates a direct

**buffer** `ByteBuffer.wrap()` - wraps an existing Java array

`clear()` prepares buffer for reading (discards previous

content) `flip()` prepares buffer for writing after it was used  
for reading

Lecture  
8 Slide  
19

# Channels

- A channel is an open connection used for reading/writing of data contained in Buffers
- Channels static class provides utility methods for

## bridging `java.io` Streams and `java.nio` Channels

- `ReadableByteChannel`, `WritableByteChannel`, and `ByteChannel` operate on `ByteBuffer`s
- `ScatteringByteChannel` and `GatheringByteChannel` operate on arrays of buffers (useful when reading fixed length heterogeneous data sequentially)

Lecture  
8 Slide  
20

## **FileChannel**

- Can be obtained with `getChannel()` either from *`FileInputStream`*, *`FileOutputStream`*, or

*RandomAccessFile* (read-only, write-only, and read write respectively)

- Provides file locking and file portion locking with `lock()` methods
- Provides memory mapping with `map()` methods
  - Memory mapped I/O maps files (whole or in parts) directly to the memory to speed up random reads/writes
- “Advanced” features depend on the support from the underlying OS

Lecture  
8 Slide  
21

## Channels and Buffers Task



- Write some more 'copy' functionality
  - With non-direct buffer
  - With direct buffer
  - With memory-mapped I/O
- Measure performance of these and the ones written before
  - Stream based
  - Buffered stream based
- Implement everything as implementing the *FileCopier* interface
- Input file, output file, and copying method should be specifiable on the command-line
- Write unit tests (you can use `File.createTempFile()` method)

# Asynchronous I/O

- `java.nio` provides APIs for asynchronous I/O
  - All other I/O in Java is synchronous (blocking)
- Channels, extending the abstract *SelectableChannel*, can work asynchronously
- *Selector* is a multiplexor, obtained using `Selector.open()`
- *SelectionKey* represents a *SelectableChannel*'s registration with a *Selector*
- *SelectableChannel* provides several `register()` methods for registering itself with a *Selector* for particular events
  - `configureBlocking()` is used for enabling asynchronous mode
- `selector.select()` is then used for waiting for registered events

