

Chapter 1: Introduction



Operating System Concepts – 10th Edition Silberschatz, Galvin and Gagne ©2018



Outline

- What Operating Systems Do
- Computer-System Organization
- Computer-System Architecture

- Operating-System Operations
 - Resource Management
 - Security and Protection
 - Virtualization
 - Distributed Systems
 - Kernel Data Structures
 - Computing Environments
-
- Free/Libre and Open-Source Operating Systems



Operating System Concepts – 10th Edition 1.2 Silberschatz, Galvin and Gagne ©2018



Objectives

- Describe the general organization of a computer system and the role of interrupts
- Describe the components in a modern, multiprocessor

computersystem

- Illustrate the transition from user mode to kernel mode
- Discuss how operating systems are used in various computing environments
- Provide examples of free and open-source operating systems



Operating System Concepts – 10th Edition 1.3 Silberschatz, Galvin and Gagne ©2018



What Does the Term Operating System Mean?

- An operating system is “fill in the blanks”
- What about:
 - Program

- Hardware
- Compiler
- Etc.



Operating System Concepts – 10th Edition 1.4 Silberschatz, Galvin and Gagne ©2018



What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware

- Operating system goals:
 - Execute user programs and make solving user problem easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner



Operating System Concepts – 10th Edition 1.5 Silberschatz, Galvin and Gagne ©2018



Computer System Structure

- Computer system can be divided into four components:
 - Hardware – provides basic computing resources
 - CPU, memory, I/O devices

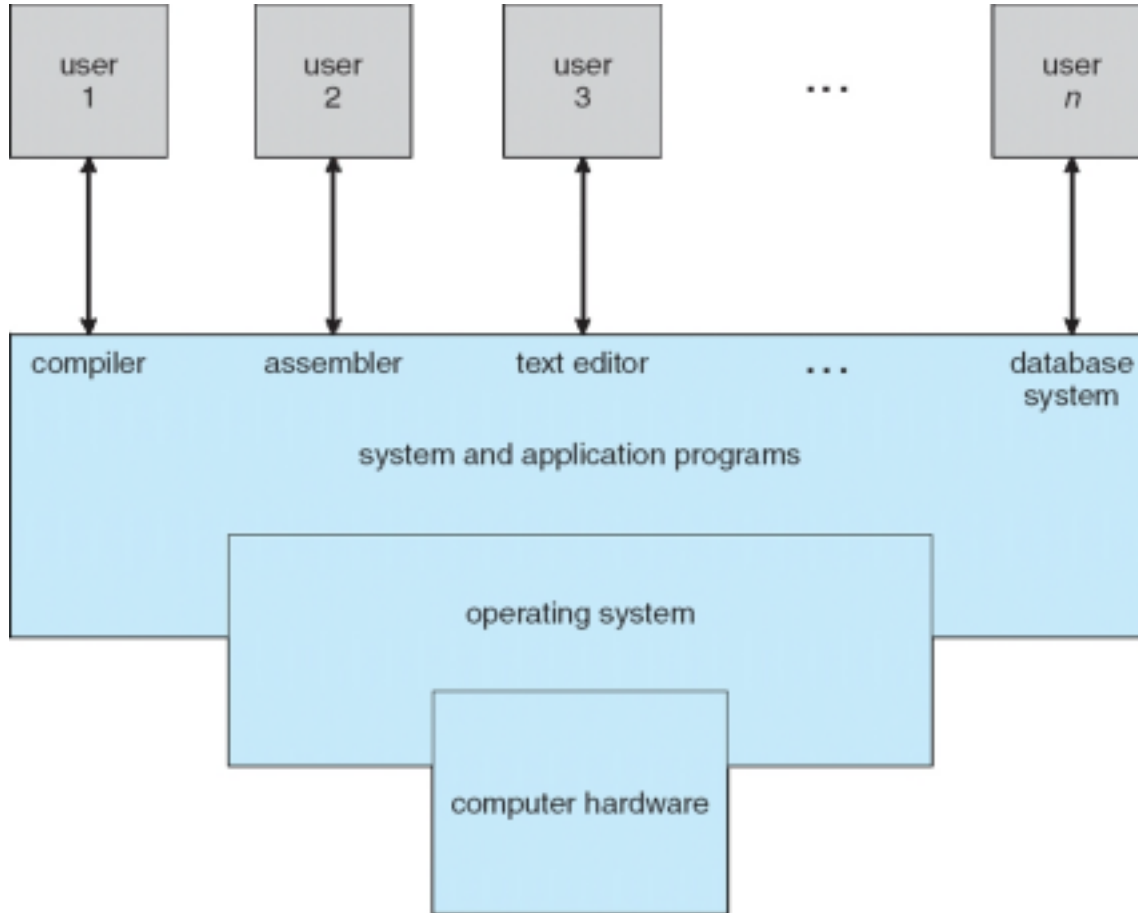
- Operating system
 - 4 Controls and coordinates use of hardware among various applications and users
- Application programs – define the ways in which the system resources are used to solve the computing problems of the users
 - 4 Word processors, compilers, web browsers, database systems, video games
- Users
 - 4 People, machines, other computers



Operating System Concepts – 10th Edition 1.6 Silberschatz, Galvin and Gagne ©2018



Abstract View of Components of Computer



Operating System Concepts – 10th Edition 1.7 Silberschatz, Galvin and Gagne ©2018



Abstract View of Components of Computer





What Operating Systems Do

- Depends on the point of view
- User's View (type of user)
 - Standalone s/m - OS is designed for **ease of use** and **good performance**
 - different terminals connected to **mainframe** or **minicomputer** - OS is designed to **maximize resource utilization**
 - users of workstation - **ease of use** and **resource utilization**
 - users of handheld devices - **ease of use** and **good performance** per amount of battery
 - optimized for usability and battery life
- Embedded systems - designed for less user interactions and **ease of use**
 - little or no user interface
 - designed to run primarily without user intervention





What Operating Systems Do (Cont.)

■ System View:

- OS is a **resource allocator**
- Manages all resources
- Decides between conflicting requests for efficient and fair resource use
- OS is a **control program** • Controls execution of programs to prevent errors and improper use of the computer






Operating System Definition

■ No universally accepted definition

- “Everything a vendor ships when you order an operating system” is a good approximation
 - But varies wildly
- “The one program running at all times on the computer” is the **kernel**, which is part of the operating system
- Everything else is either
 - A **system program** (ships with the operating system, but not part of the kernel) , or
 - An **application program**, all programs not associated with the operating system
 - Today’s OSes for general purpose and mobile computing also include **middleware** – a set of software frameworks that provide additional services to application developers such as databases, multimedia, graphics etc.



Operating System Concepts – 10th Edition 1.11 Silberschatz, Galvin and Gagne ©2018





Computer System Organization.

Computer-system operation

- One or more CPUs, device controllers connect through common **bus** providing access to shared memory
- Concurrent execution of CPUs and devices competing for memory cycles





Operating System Concepts – 10th Edition 1.13 Silberschatz, Galvin and Gagne ©2018



Computer-System Operation

- I/O devices and the CPU can execute concurrently
- Each device controller is in charge of a particular device type
- Each device controller has a local buffer

- Each device controller type has an operating system **device driver** to manage it
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller ▪ Device controller informs CPU that it has finished its operation by causing an **interrupt**



Operating System Concepts – 10th Edition 1.14 Silberschatz, Galvin and Gagne ©2018



Computer Startup

- **Bootstrap program** is loaded at power-up or reboot ▪ Typically stored in ROM or EPROM, generally known as **firmware**
 - Initializes all aspects of system: CPU registers, memory, device

controllers and other initial setups

- Loads operating system kernel and OS starts with the execution of 'init' process
- **Bootstrap program**
 - Initializes regs., memory, I/O devices
 - locates and loads kernel into memory
 - starts with 'init' process
 - waits for interrupt from user



Operating System Concepts – 10th Edition 1.15 Silberschatz, Galvin and Gagne ©2018



Interrupt Handling ▪ Occurrence of an event is signalled by **interrupt**

- **Hardware interrupt** - triggered by an event external to the processor

- **Software interrupt** - triggered by executing special instruction/operation (API or system calls)
- Interrupt transfers control to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines
- Interrupt architecture must save the address of the interrupted instruction
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request
- An operating system is **interrupt driven**



Operating System Concepts – 10th Edition 1.16 Silberschatz, Galvin and Gagne ©2018



Interrupt Driven Program Execution





Interrupt Timeline





Interrupt Handling

- The operating system preserves the state of the CPU by storing the registers and the program counter
- Determines which type of interrupt has occurred:
- Separate segments of code determine what action should be taken for each type of interrupt






Interrupt-drive/OCycle



Operating System Concepts – 10th Edition 1.20 Silberschatz, Galvin and Gagne ©2018





Storage Structure

- Main memory – only large storage media that the CPU can access directly
 - Typically, **volatile**
 - Typically, **random-access memory** in the form of **Dynamic Random-access Memory (DRAM)**
- Secondary storage – extension of main **memory** that provides large **nonvolatile** storage capacity





Storage Structure(Cont.)

- **Hard Disk Drives (HDD)** – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into **tracks**, which are subdivided into **sectors**
 - The **disk controller** determines the logical interaction between the device and the computer

- **Non-volatile memory (NVM)** devices– faster than harddisks, nonvolatile
 - Various technologies
 - Becoming more popular as capacity and performance increases, price drops



Operating System Concepts – 10th Edition 1.23 Silberschatz, Galvin and Gagne ©2018



Storage Definitions and Notation Review

The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a

computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers, it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes. A **kilobyte**, or KB, is 1,024 bytes; a **megabyte**, or **MB**, is $1,024^2$ bytes; a **gigabyte**, or GB, is $1,024^3$ bytes; a **terabyte**, or **TB**, is $1,024^4$ bytes; and a **petabyte**, or **PB**, is $1,024^5$ bytes. Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).



Operating System Concepts – 10th Edition 1.24 Silberschatz, Galvin and Gagne ©2018



Storage Hierarchy

- Storage systems organized in hierarchy

- Speed
 - Cost
 - Volatility
- **Caching** – copying information into faster storage system; main memory can be viewed as a cache for secondary storage
 - **Device Driver** for each device controller to manage I/O • Provides uniform interface between controller and kernel



Operating System Concepts – 10th Edition 1.25 Silberschatz, Galvin and Gagne ©2018



Storage-Device Hierarchy



Operating System Concepts – 10th Edition 1.26 Silberschatz, Galvin and Gagne ©2018



How a Modern Computer Works



A von Neumann architecture



Operating System Concepts – 10th Edition 1.27 Silberschatz, Galvin and Gagne ©2018



I/O Structure

- Two methods for handling I/O
 - After I/O starts, control returns to user program only upon I/O completion
 - After I/O starts, control returns to user program without waiting for I/O completion



- After I/O starts, control returns to user program only upon I/O completion
 - Wait instruction idles the CPU until the next interrupt
 - Wait loop (contention for memory access)
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing
- After I/O starts, control returns to user program without waiting for I/O completion
 - **System call** – request to the OS to allow user to wait for I/O completion
 - **Device-status table** contains entry for each I/O device indicating its type, address, and state
 - OS indexes into I/O device table to determine device status and to modify table entry to include interrupt



Direct Memory Access Structure

- Used for high-speed I/O devices able to transmit information at close to memory speeds
- Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention
- Only one interrupt is generated per block, rather than the one interrupt per byte



Operating System Concepts – 10th Edition 1.30 Silberschatz, Galvin and Gagne ©2018



Computer-System Architecture

- Most systems use a single general-purpose processor • Most systems have special-purpose processors as well
- **Multiprocessors** systems growing in use and importance • Also known as **parallel systems**, **tightly-coupled systems** •

Advantages include:

1. **Increased throughput** : No. programs executed per unit time is more; increasing number of processors do not always guarantee increased performance due to the overhead of complexity and cost
2. **Economy of scale** – costs less than equivalent no. of many single processors due to sharing of peripherals, mass storage and power supplies
3. **Increased reliability** – if one processor fails, the system is not halted, it only slows down; the job of the failed processor is taken up by other processor

- graceful degradation or fault tolerant

Operating System Concepts – 10th Edition 1.31 Silberschatz, Galvin and Gagne ©2018



Symmetric Multiprocessing Architecture

Two types of multiprocessor systems:

1. Asymmetric Multiprocessing

(Master/Slave) – each processor is assigned a specific task by master processor. It schedules and allocates work to slave processors

2. Symmetric Multiprocessing (SMP) –

each processor performs all tasks; all processors are considered peers; Windows, MacOS, Linux support SMP





A Dual-Core Design

- Multiple processing units (cores) fabricated on a single chip - **multicore processors**
 - The term processor is then used for the complete chip
- The communication b/w processors within a chip is more faster than communication between two single processors





Clustered Systems

- Like multiprocessor systems, but multiple systems working together and sharing software resources
 - Usually sharing storage via a **storage-area network (SAN)** - where resources can be shared with dozens of systems in a cluster, that are separated by miles.
 - Provides a **high-availability** of resources and services which survive failures; the service continues even if one or more systems in the cluster fail; duplication of s/w resources
 - 4 **Asymmetric clustering** has one machine in hot-standby mode
 - 4 **Symmetric clustering** has multiple nodes running applications, monitoring each other
 - Some clusters are for **high-performance computing (HPC)** 4
 - Applications must be written to use **parallelization**
 - Parallel clusters allow multiple hosts to access

samedataonthesharedstorage



Operating System Concepts – 10th Edition 1.34 Silberschatz, Galvin and Gagne ©2018



Clustered Systems





Operating System Concepts – 10th Edition 1.35 Silberschatz, Galvin and Gagne ©2018




Operating System Structure

- **Multiprogramming** (**Batch system**) needed for efficiency • Single user cannot keep CPU and I/O devices busy at all times • Multiprogramming organizes jobs (code and data) so CPU always has one to execute

- A subset of total jobs in system is kept in memory • One job selected and run via **job scheduling** • When it has to wait (for I/O for example), OS switches to and executes another job
- Multiprogrammed s/ms provide environment in which various s/m resources are utilized effectively, but they do not provide user interaction with the computer system



Operating System Concepts – 10th Edition 1.36 Silberschatz, Galvin and Gagne ©2018



Memory Layout for

MultiprogrammedSystem



Operating System Concepts – 10th Edition 1.37 Silberschatz, Galvin and Gagne ©2018



Operating SystemStructure

- **Timesharing (multitasking)** is logical extension in which a single CPU executes multiple jobs by switching among them

switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing

- **Response time** should be < 1 second
- Each user has at least one program executing in memory □ **process** • If several jobs ready to run at the same time □ **CPU scheduling**
- If processes don't fit in memory, **swapping** moves them in and out to run • **Virtual memory** allows execution of processes not completely in memory





Operating-SystemOperations

- **Interrupt driven** (hardware and software)
 - Hardware interrupt by one of the devices
 - Software interrupt (**exception** or **trap**):
 - 4 Software error (e.g., division by zero)
 - 4 Request for operating system service
 - 4 Other process problems include infinite loop, processes modifying each other or the operating system





Dual-mode Operation

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
- **Mode bit** provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code.
 - When a user is running □ mode bit is “user” (set to 1) • When kernel code is executing □ mode bit is “kernel” (set to 0) ■ Some instructions designated as **privileged**, only executable in kernel mode.
- Eg: changing the mode bit; changing the priority level etc.





Dual-mode Operation(Cont.)

- How do we guarantee that user does not explicitly set the mode bit to “kernel”?
- When the system starts executing it is in kernel mode ■ When control is given to a user program the mode-bit changes to “user mode”.
- When a user issues a system call it results in an interrupt, which traps to the operating system. At that time, the mode-bit is set to “kernel mode”.





Transition from User to Kernel Mode





Timer

- Operating system uses timer to control the CPU. A user program cannot hold CPU for a long time, this is prevented with the help of timer.
- Timer to prevent infinite loop (or process hogging resources) • Timer is set to interrupt the computer after some time period □ Fixed timer - After a fixed time, the process under execution is interrupted. • Variable timer - Interrupt occurs after varying interval. • Fixed clock and a counter that is decremented every time the clock ticks
 - Operating system set the counter (privileged instruction) • When counter zero generate an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time





Process Management

A process is a program in execution. It is a unit of work within the system. Program is a ***passive entity***; process is an ***active entity***.
▪ Process needs resources to accomplish its task
• CPU, memory, I/O, files

- Initialization data
 - Process termination requires reclaim of any reusable resources
- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one program counter per thread
 - Typically, system has many processes, some user, some operating system running concurrently on one or more CPUs
 - Concurrency by multiplexing the CPUs among the processes/ threads





Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- Creating and deleting both user and system processes
- Suspending and resuming processes
- Providing mechanisms for process synchronization
- Providing mechanisms for process communication
- Providing mechanisms for deadlock handling





Memory Management

- To execute a program all (or part) of the instructions must be in memory
- All (or part) of the data that is needed by the program must be in memory
 - To improve both the utilization of the CPU several programs are kept in memory, creating a need for memory management.
- Memory management determines what is in memory and when
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and de-allocating memory space as needed





File-systemManagement

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit - **file**
 - Different types of physical media - Magnetic disk, optical disk, andmagnetic tape`
 - Each medium is controlled by device (i.e., disk drive, tapedrive)
 - 4 Varying properties include access speed, capacity, data-transferrate, access method (sequential or random)
- File-System management
 - Files usually organized into directories
 - Multiple users - Access control on most systems todeterminewhocanaccess what
 - OS activities include
 - 4 Creating and deleting files and directories
 - 4 Supporting primitives to manipulate files and directories
 - 4 Mapping files onto secondary storage
 - 4 Backup files onto stable (non-volatile) storage media





Mass-Storage Management

- Main memory is small and volatile – secondary storage backup main memory
- Usually, disks used to store data that does not fit in main memory or data that must be kept for a “long” period of time
- Proper management is of central importance
- Entire speed of computer operation hinges on disk subsystem and its algorithms
- OS activities
 - Mounting and unmounting
 - Free-space management
 - Storage allocation
 - Disk scheduling
 - Partitioning
 - Protection





Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
- Cache management important design problem
- Cache size and replacement policy





Characteristics of Various Types of Storage



Movement between levels of storage hierarchy can be explicit or implicit





Migration of data “A” from Disk to Register

- Multitasking environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



- Multiprocessor environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache
- Distributed environment situation even more complex
 - Several copies of a datum can exist
 - Various solutions covered in Chapter 19





I/O Subsystem

- One purpose of OS is to hide peculiarities of hardware devices from the user
- I/O subsystem responsible for
 - Memory management of I/O including buffering (storing data temporarily while it is being transferred), caching (storing parts of data in faster storage for performance), spooling (the overlapping of output of one job with input of other jobs)
 - General device-driver interface
 - Drivers for specific hardware devices



Protection and Security

- **Protection** – mechanism for controlling access of processes or users to resources defined by the OS
- **Security** – defense of the system against internal and external attacks
 - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service



Protection

- Systems generally first distinguish among users, to determine who can do what
 - User identities (**user IDs**, security IDs) include name and associated number, one per user

- User ID then associated with all files, processes of that user to determine access control
- Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
- **Privilege escalation** allows user to change to effective ID with more rights



Operating System Concepts – 10th Edition 1.54 Silberschatz, Galvin and Gagne ©2018



Computing Environments- Traditional

- Stand-alone general purpose machines
- But blurred as most systems interconnect with others (i.e., the Internet)
- **Portals** provide web access to internal systems
- **Network computers** (thin

clients) are like Web terminals ■ Mobile computers interconnect via **wireless networks** ■ Networking becoming ubiquitous – even home systems use **firewalls** to protect home computers from Internet attacks



Operating System Concepts – 10th Edition 1.55 Silberschatz, Galvin and Gagne ©2018



Computing Environments- Mobile

- Handheld smartphones, tablets, etc
- What is the functional difference between them and a “traditional” laptop?
- Extra feature – more OS features (GPS, gyroscope) ■ Allows new types of apps like **augmented reality** ■ Use IEEE 802.11 wireless, or cellular data

networks for connectivity

- Leaders are **Apple iOS** and **Google Android**



Operating System Concepts – 10th Edition 1.56 Silberschatz, Galvin and Gagne ©2018



Computing Environments—Distributed.

Distributed computing

- Collection of separate, possibly heterogeneous, systems networked together
 - 4 **Network** is a communications path, **TCP/IP** most common
 - **Local Area Network (LAN)**

- **Wide Area Network (WAN)**
- **Metropolitan Area Network (MAN)**
- **Personal Area Network (PAN)**
- **Network Operating System** provides features between systems across network
 - 4 Communication scheme allows systems to exchange messages
 - 4 Illusion of a single system



Operating System Concepts – 10th Edition 1.57 Silberschatz, Galvin and Gagne ©2018



Computing Environments–Client-Server

■ Client-Server Computing

- Dumb terminals supplanted by smart PCs
- Many systems now **servers**, responding to requests generated by **clients**
 - 4 **Compute-server system** provides an interface to client to request

services (i.e., database)

- 4 **File-server system** provides interface for clients to store and retrieve files



Operating System Concepts – 10th Edition 1.58 Silberschatz, Galvin and Gagne ©2018



Computing Environments- Peer-to-Peer

- Another model of distributed system
- P2P does not distinguish clients and servers
- Instead all nodes are considered peers
- May each act as client, server or both

- Node must join P2P network
 - 4 Registers its service with central lookup service on network, or
 - 4 Broadcast request for service and respond to requests for service via

discovery protocol

- Examples include Napster and Gnutella, **Voice over IP (VoIP)** such as Skype



Operating System Concepts – 10th Edition 1.59 Silberschatz, Galvin and Gagne ©2018



Computing Environments – Real-Time Embedded Systems

- Real-time embedded systems most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, **real-time OS**
 - Use expanding

- Many other special computing environments as well • Some have OSes, some perform tasks without an OS
- Real-time OS has well-defined fixed time constraints • Processing **must** be done within constraint
 - Correct operation only if constraints met



Operating System Concepts – 10th Edition 1.60 Silberschatz, Galvin and Gagne ©2018



Computing Environments- Virtualization

- Allows operating systems to run applications within other OSes • Vast and growing industry
- **Emulation** - creates an environment that imitates the properties of one system onto another, mimics the properties of one processor to run in another platform

efficiently

- hardware component is replaced by software based construct
- used when source CPU type different from target type
 - run programs designed for a completely different architecture on an x86 PC
 - running an old game designed for obsolete platforms on today's modern systems (i.e. PowerPC to Intel x86)
- Generally slowest method
- When computer language not compiled to native code—**Interpretation**



Operating System Concepts – 10th Edition 1.61 Silberschatz, Galvin and Gagne ©2018



Computing Environments- Virtualization

- **Virtualization** – OS natively compiled for CPU, running **guest** OSes also natively compiled
 - Consider VMware running WinXP guests, each running applications, all on native WinXP **host** OS

- **VMM** (virtual machine Manager) provides virtualizationservices▪ Use cases involve laptops and desktops running multipleOSesforexploration or compatibility
- Apple laptop running Mac OS X host, Windows as a guest • Developing apps for multiple OSes without having multiplesystems• Executing and managing compute environments withindatacenters▪ VMM can run natively, in which case they are also the host



Operating System Concepts – 10th Edition 1.62 Silberschatz, GalvinandGagne©2018



Computing Environments- Virtualization





Computing Environments- Virtualization

- Emulation requires software bridge ■ With virtualization hardware can be directly accessed whereas in Emulator the OS does not run on the physical hardware
- Emulator requires interpreter to translate the source code to host's readable format, to further process it.
- Emulators are slower than VMs
- Emulators do not rely on CPU while VMs make use of CPU ■ VM solution is costlier and more complex than Emulation technique
- VM provides more throughput, minimal overhead with a better backup and recovery solution





Computing Environments – CloudComputing

Delivers computing, storage, even apps as a service across a network. Logical extension of virtualization because it uses virtualization as the base for its functionality.

- Amazon **EC2** has thousands of servers, millions of virtual machines, petabytes of storage available across the Internet, pay based on usage. Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., word processor)
 - Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)





Computing Environments – CloudComputing

- Cloud computing environments composed of traditional OSES, plus VMMs, plus cloud management tools
 - Internet connectivity requires security like firewalls
 - Load balancers spread traffic across multiple applications





Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary **closed-source**
- Counter to the **copy protection** and **Digital Rights Management (DRM)** movement
- Started by **Free Software Foundation (FSF)**, which has “copyleft” **GNU Public License (GPL)**
- Examples include **GNU/Linux** and **BSD UNIX** (including core of **Mac OS X**), and many more
 - Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - <http://www.virtualbox.com>)
 - Use to run guest operating systems for exploration



Operating System Concepts – 10th Edition 1.67 Silberschatz, Galvin and Gagne ©2018



The Study of Operating Systems

There has never been a more interesting time to study operating systems, and it has never been easier. The open-source movement has overtaken operating systems, causing many of them to be made available in both source and binary (executable) format. The list of operating systems available in both formats includes Linux, BSD UNIX, Solaris, and part of macOS. The availability of source code allows us to study operating systems from the inside out. Questions that we could once answer only by looking at documentation or the behavior of an operating system we can now answer by examining the code itself.

Operating systems that are no longer commercially viable have been open-sourced as well, enabling us to study how systems operated in a time of fewer CPU, memory, and storage resources. An extensive but incomplete list of open-source operating-system projects is available from https://curlie.org/Computers/Software/Operating_Systems/Open_Source/

In addition, the rise of virtualization as a mainstream (and frequently free) computer function makes it possible to run many operating systems on top of one core system. For example, VMware (<http://www.vmware.com>) provides a free “player” for Windows on which hundreds of free “virtual appliances” can run. Virtualbox (<http://www.virtualbox.com>) provides a free, open-source virtual machine manager on many operating systems. Using such tools, students can try out hundreds of operating systems without dedicated hardware.

The advent of open-source operating systems has also made it easier to make the move from student to operating-system developer. With some knowledge, some effort, and an Internet connection, a student can even create a new operating-system distribution. Just a few years ago, it was difficult or impossible to get access to source code. Now, such access is limited only by how much interest, time, and disk space a student has.



Operating System Concepts – 10th Edition 1.68 Silberschatz, Galvin and Gagne ©2018

Chapter 2a:

Operating-System Services





Outline

- Operating System Services
- User and Operating System-Interface
- System Calls
- System Services
- Linkers and Loaders
- Why Applications are Operating System Specific



Objectives

- Identify services provided by an operating system
- Illustrate how system calls are used to provide operating system services



Operating System Services

- Operating systems provide an environment for execution of programs and services to programs and users

- One set of operating-system services provides functions that are helpful to the user:
 - **User interface** - Almost all operating systems have a user interface (**UI**).
 - 4 Varies between **Command-Line (CLI)**, **Graphics User Interface (GUI)**, **touch-screen**, **Batch**
 - **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
 - **I/O operations** - A running program may require I/O, which may involve a file or an I/O device



Operating System Concepts – 10th Edition 1.72 Silberschatz, Galvin and Gagne ©2018



Operating System Services (Cont.)

- One set of operating-system services provides functions that are helpful to the user (Cont.):

- **File-system manipulation** - The file system is of particular interest. Programs need to read and write files and directories, create and delete them, search them, list file information, permission management.
- **Communications** – Processes may exchange information, on the same computer or between computers over a network
 - 4 Communications may be via shared memory or through message passing (packets moved by the OS)



Operating System Concepts – 10th Edition 1.73 Silberschatz, Galvin and Gagne ©2018



Operating System Services (Cont.)

- One set of operating-system services provides functions that are helpful to the

user (Cont.):

- **Error detection** – OS needs to be constantly aware of possible errors
 - 4 May occur in the CPU and memory hardware, in I/O devices, in user program
 - 4 For each type of error, OS should take the appropriate action to ensure correct and consistent computing
 - 4 Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system





Operating System Services(Cont.)

- Another set of OS function exists for ensuring the efficient operation of the system itself via resource sharing
 - **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
 - 4 Many types of resources - CPU cycles, main memory, file storage, I/O devices.
 - **Logging** - To keep track of which users use how much and what kinds of computer resources
 - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
 - 4 **Protection** involves ensuring that all access to system resources is controlled
 - 4 **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts



Operating System Concepts – 10th Edition 1.75 Silberschatz, Galvin and Gagne ©2018



A View of Operating System Services





User Operating System Interface

- CLI -- command line interpreter
 - allows direct command entry
- GUI – graphical user interface
- Touchscreen Interfaces
- Batch



CLI

- Sometimes implemented in kernel, sometimes by system program
- Sometimes multiple flavors implemented – **shells** ▪ Primarily fetches a command from user and executes it
- Sometimes commands built-in, sometimes just names of programs

- If the latter, adding new features doesn't require shell modification



Operating System Concepts – 10th Edition 1.78 Silberschatz, Galvin and Gagne ©2018



Bourne Shell Command Interpreter



- User-friendly **desktop** metaphor interface
 - Usually mouse, keyboard, and monitor
 - **Icons** represent files, programs, actions, etc.
 - Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**))
 - Invented at Xerox PARC
- Many systems now include both CLI and GUI interfaces
 - Microsoft Windows is GUI with CLI “command” shell
 - Apple Mac OS X is “Aqua” GUI interface with UNIX kernel underneath and shells available
 - Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)

