

Chapter 10: Mass-StorageSystems



Operating System Concepts – 9th Edition Silberschatz, Galvin and Gagne ©2013



Chapter 10: Mass-StorageSystems

- Overview of Mass Storage Structure

- Disk Structure
- Disk Attachment
- Disk Scheduling



Operating System Concepts – 9th Edition 10.2 Silberschatz, Galvin and Gagne ©2013



Objectives

- To describe the physical structure of secondary storage devices and its effects on the uses of the devices
- To explain the performance characteristics of mass-storage devices
- To evaluate disk scheduling algorithms





Moving-headDiskMechanism





Hard Disk Performance

- **Access Latency = Average access time** = average seektime + average latency
 - For fastest disk 3ms + 2ms = 5ms
 - For slow disk 9ms + 5.56ms = 14.56ms
- Average I/O time = average access time + (amount to transfer / transfer rate) + controller overhead
- For example to transfer a 4KB block on a 7200 RPM disk with a 5ms average seek time, 1Gb/sec transfer rate with a .1ms controller overhead =
 - 5ms + 4.17ms + 0.1ms + transfer time =
 - Transfer time = $4\text{KB} / 1\text{Gb/s} * 8\text{Gb} / \text{GB} * 1\text{GB} / 1024^2\text{KB} = 32 / (1024^2) = 0.031 \text{ ms}$
 - Average I/O time for 4KB block = 9.27ms + .031ms = 9.301ms





Disk Structure

- Disk drives are addressed as large 1-dimensional arrays of **logical blocks**, where the logical block is the smallest unit of transfer
 - Low-level formatting creates **logical blocks** on physical media
- The 1-dimensional array of logical blocks is mapped into the sectors of the disk sequentially
 - Sector 0 is the first sector of the first track on the outermost cylinder
 - Mapping proceeds in order through that track, then the rest of the tracks in that cylinder, and then through the rest of the cylinders from outermost to innermost
 - Logical to physical address should be easy ⁴ Except for bad sectors
 - ⁴ Non-constant # of sectors per track via constant angular velocity





Disk Scheduling

- The operating system is responsible for using hardware efficiently — for the disk drives, this means having a fast access time and large disk bandwidth
- Access time has two components:
 - Seek time and Rotational latency
- Seek time: Time for the disk arm to move the heads to the cylinder containing the desired sector.
- Rotational latency : additional time for the disk to rotate the desired sector to the disk head.
- Minimize seek time
- Disk **bandwidth** is the total number of bytes transferred, divided by the total time between the first request for service and the completion of the last transfer





Disk Scheduling(Cont.)

There are many sources of disk I/O request

- System processes
- OS
- Users processes
- I/O request includes input or output mode, disk address, memory address, number of sectors to transfer
 - OS maintains queue of requests, per disk or device
- Idle disk can immediately work on I/O request, busy disk means work must queue
 - Optimization algorithms only make sense when a queue exists





Disk Scheduling(Cont.)

- Note that drive controllers have small buffers and can manage a queue of I/O requests (of varying “depth”)
- Several algorithms exist to schedule the servicing of disk I/O requests ■ The analysis is true for one or many platters
- Suppose a disk is having 200 cylinders numbered from 0 to 199. The disk is currently servicing at cylinder 53 and previous request was at cylinder 60. The Queue of pending requests in FIFO order is -

98, 183, 37, 122, 14, 124, 65, 67

Calculate the total distance the read/write head will travel.

FCFS

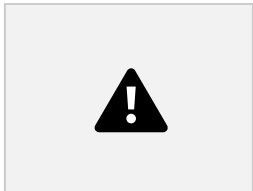
Illustration shows total head movement of 640 cylinders:

$$640 = (98 - 53) + (183 - 98) + (183 - 37) + (122 - 37) + (122 - 14) + (124 - 14) + (124 - 65) + (67 - 65)$$



Operating

System Concepts – 9th Edition 10.10 Silberschatz, Galvin and Gagne ©2013



SSTF

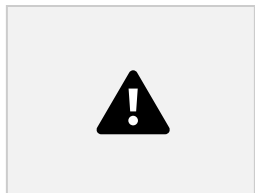
- Shortest Seek Time First selects the request with the minimum seek time from the current head position: SSTF scheduling is a form of SJF scheduling; may cause starvation of some requests

- Illustration shows total head movement of 236 cylinders

$$236 = (65-53) + (67-65) + (67-37) + (37-14) + (98-14) + (122-98) + (124-122) + (183-124)$$



Operating System Concepts – 9th Edition 10.11 Silberschatz, Galvin and Gagne ©2013



SCAN

- The disk arm starts at one end of the disk, and moves toward the other end, servicing requests until it gets to the other end of the disk, where the head

movement is reversed and servicing continues.

- **SCAN algorithm** Sometimes called the **elevator algorithm** ■ Illustration shows total head movement of 236 cylinders
- But note that if requests are uniformly dense, largest density at other end of disk and those wait the longest



Operating System Concepts – 9th Edition 10.12 Silberschatz, Galvin and Gagne ©2013



SCAN(Cont.) 236 =
(53-37)+(37-14)+(14-0)+(65-0)+(67-65)+(98-67)+(122-98)+(124-122)+(18

(OR) Total Head Movements = $(53 - 0) + (183 - 0) = 236$



Operating System

Concepts – 9th Edition 10.13 Silberschatz, Galvin and Gagne ©2013



C-SCAN

- Provides a more uniform wait time than SCAN
- The head moves from one end of the disk to the other(increasing order of cylinder number), servicing requests as it goes
 - When it reaches the other end, however, it immediately returns to the beginning of the disk, without servicing any requests on the return trip
- Treats the cylinders as a circular list that wraps around from the last cylinder to the first one
- Total number of cylinders?



Operating System Concepts – 9th Edition 10.14 Silberschatz, Galvin and Gagne ©2013



C-SCAN(Cont.) Total Head Movements = $(199 - 53) +$

$$(199-0)+(37-0) = 382$$



Operating System Concepts – 9th Edition 10.15 Silberschatz, Galvin and Gagne ©2013



C-LOOK

■ LOOK a version of SCAN, C-LOOK a version of C-SCAN

- Arm only goes as far as the last request in each direction, then reverses direction immediately, without first going all the way to the end of the disk
- Total number of cylinders?







$$\text{Head Movements} = (183 - 53) + (183 - 14) + (37 - 14) = 322$$



Selecting a Disk-Scheduling Algorithm

- SSTF is common and has a natural appeal
- SCAN and C-SCAN perform better for systems that place a heavy load on the disk
 - Less starvation
- Performance depends on the number and types of requests ■ Requests for disk service can be influenced by the file-allocation method
- The disk-scheduling algorithm should be written as a separate module of the operating system, allowing it to be replaced with a different algorithm if necessary
- Either SSTF or LOOK is a reasonable choice for the default algorithm



PROTECTION:

■ GOALS OF PROTECTION

- Protection is a mechanism for **controlling the access** of programs, processes, or users to the resources defined by a computer system. Protection ensures that only processes that have gained proper authorization from the operating system can operate on the files, memory segments, CPU, and other resources of a system.
- Protection is required to prevent mischievous, intentional violation of an access restriction by a user.

■ PRINCIPLES OF PROTECTION

- A key, time-tested guiding principle for protection is the 'principle of least privilege'. It dictates that programs, users, and even systems be given just enough privileges to perform their tasks. An operating system provides mechanisms to enable privileges when they are needed and to disable them when they are not needed.



Operating System Concepts – 9th Edition 10.20 Silberschatz, Galvin and Gagne ©2013



PROTECTION:

■ DOMAIN OF PROTECTION

- A computer system is a collection of processes and objects. *Objects* are both **hardware objects** (such as the CPU, memory segments, printers, disks, and tape drives) and **software objects** (such as files, programs, and semaphores). Each object (resource) has a unique name that differentiates it from all other objects in the system.
- The operations that are possible may depend on the object. For example, a CPU can only be executed on. Memory segments can

be read and written, whereas a CD-ROM or DVD-ROM can only be read. Tape drives can be read, written, and rewound. Data files can be created, opened, read, written, closed, and deleted; program files can be read, written, executed, and deleted.

- A process should be allowed to access only those resources
 - a) for which it has authorization
 - b) currently requires to complete process

