

Chapter 10: Virtual Memory



Operating System Concepts – 10th Edition Silberschatz, Galvin and Gagne ©2018



Outline

- Background

- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocating Kernel Memory
- Other Considerations
- Operating-System Examples



Operating System Concepts – 10th Edition 10.2 Silberschatz, Galvin and Gagne ©2018



Objectives

- Define virtual memory and describe its benefits.
- Illustrate how pages are loaded into memory using demand paging. ▪ Apply the FIFO, optimal, and LRU page-replacement algorithms. ▪ Describe the working set of a process and explain how it is related to program locality.
- Describe how Linux, Windows 10, and Solaris manage virtual memory.
- Design a virtual memory manager simulation in the C programming language.



Operating System Concepts – 10th Edition 10.3 Silberschatz, Galvin and Gagne ©2018



Background

- Code needs to be in memory to execute, but entire program rarely used

- Error code, unusual routines, large data structures
- Entire program code not needed at same time
- Consider ability to execute partially-loaded program
- Program no longer constrained by limits of physical memory
- Each program takes less memory while running -> more programs run at the same time
 - 4 Increased CPU utilization and throughput with no increase in response time or turnaround time
- Less I/O needed to load or swap programs into memory -> each user program runs faster



Operating System Concepts – 10th Edition 10.4 Silberschatz, Galvin and Gagne ©2018



Virtual memory

- **Virtual memory** – separation of user logical memory from physical memory

- Only part of the program needs to be in memory for execution
- Logical address space can therefore be much larger than physical address space
 - Allows address spaces to be shared by several processes
- Allows for more efficient process creation
- More programs running concurrently
- Less I/O needed to load or swap processes



Operating System Concepts – 10th Edition 10.5 Silberschatz, Galvin and Gagne ©2018



Virtual memory(Cont.)

- **Virtual address space** – logical view of how process is stored in memory


- Usually start at address 0, contiguous addresses until end of space
 - Meanwhile, physical memory organized in page frames
 - MMU must map logical to physical
- Virtual memory can be implemented via:
- Demand paging
 - Demand segmentation



Operating System Concepts – 10th Edition 10.6 Silberschatz, Galvin and Gagne ©2018



Virtual Memory That is Larger



ThanPhysicalMemory



Operating System Concepts – 10th Edition 10.7 Silberschatz, Galvin and Gagne ©2018




Virtual-addressSpace

- Usually design logical address space for the stack to start at Max logical address and grow “down” while heap grows “up”

- Maximizes address space use
- Unused address space between two is hole

4 No physical memory needed until heap or grows to a given new


- Enables **sparse** address spaces with holes left for growth, dynamically linked libraries, etc.
- System libraries shared via mapping virtual address space
 - Shared memory by mapping pages read- write into virtual address space
 - Pages can be shared during **fork()** , speeding process creation



the

stack
page

into






Shared Library Using Virtual Memory





DemandPaging

- Could bring entire process into memory at load time
- Or bring a page into memory only when it is needed
 - Less I/O needed, no unnecessary I/O
 - Less memory needed
 - Faster response
 - More users
- Similar to paging system with (diagram on right)
- invalid reference \Rightarrow abort
 - Not-in-memory \Rightarrow bring to
- **Lazy swapper** – never swaps a memory unless page will be
 - Swapper that deals with pages is a **pager**



swapping

memory
page into
needed




Basic Concepts

- With swapping, the pager guesses which pages will be used before swapping them out again
- How to determine that set of pages?
- Need new MMU functionality to implement demand paging
- If pages needed are already **memory resident**
 - No difference from non demand-paging
 - If page needed and not memory resident
 - Need to detect and load the page into memory from storage
- ⁴ Without changing program behavior
 - Without programmer needing to change code
 - Use page table with valid-invalid bit (see chapter 9)



Page table with Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated (**v** ⇒ in-memory, **i** ⇒ not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries
- Example of a page table snapshot:





- During MMU address translation, if valid–invalid bit in the page table entry is **i**
⇒ page fault



Operating System Concepts – 10th Edition 10.12 Silberschatz, Galvin and Gagne ©2018


Page Table When Some Pages Are Not in Main Memory





Operating

System Concepts – 10th Edition 10.13 Silberschatz, Galvin and Gagne ©2018



Steps in Handling Page Fault

1. If there is a reference to a page, first reference to that page will trap to operating system • Page fault
2. Operating system looks at another table to decide: • Invalid reference \Rightarrow abort • Just not in memory (go to step 3)
3. Find free frame (what if there is none?)
4. Swap page into frame via scheduled disk operation
5. Reset tables to indicate page now in memory Set validation bit = **v**
6. Restart the instruction that caused the page fault



Operating System Concepts – 10th Edition 10.14 Silberschatz, Galvin and Gagne ©2018



Steps in Handling a Page Fault(Cont.)



Operating System Concepts – 10th Edition 10.15 Silberschatz, Galvin and Gagne ©2018



Aspects of Demand Paging

■ Pure demand paging: start process with *no* pages in memory • OS sets instruction pointer to first

instruction of process, non-memory-resident -> page fault

- And for every other process pages on first access ■ Actually, a given instruction could access multiple pages->multiple page faults
- Consider fetch and decode of instruction which adds 2 numbers from memory and stores result back to memory ■ Hardware support needed for demand paging
- Page table with valid / invalid bit
- Secondary memory (swap device with **swap space**)
- Instruction restart

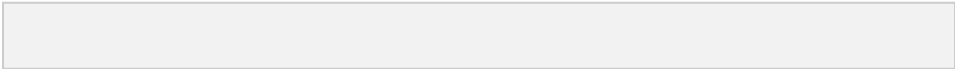


Operating System Concepts – 10th Edition 10.16 Silberschatz, Galvin and Gagne ©2018



Free-FrameList

- When a page fault occurs, the operating system must bring the desired page from secondary storage into main memory.
- Most operating systems maintain a **free-frame list** -- a pool of free frames for satisfying such requests.

- 
- Operating system typically allocate free frames using a technique known as **zero-fill-on-demand** -- the content of the frames zeroed-out before being allocated.
 - When a system starts up, all available memory is placed on the free-frame list.



Operating System Concepts – 10th Edition 10.17 Silberschatz, Galvin and Gagne ©2018



Stages in Demand Paging–Worse Case 1. Trap

to the operating system 2. Save the user registers and process state 3.

Determine that the interrupt was a page fault 4. Check that the page reference was legal and determine the location of the page on the disk

5. Issue a read from the disk to a free frame:

- a) Wait in a queue for this device until the read request is serviced
- b) Wait for the device seek and/or latency time
- c) Begin the transfer of the page to a free frame



Operating System Concepts – 10th Edition 10.18 Silberschatz, Galvin and Gagne ©2018



Stages in Demand Paging(Cont.)

- 6. While waiting, allocate the CPU to some other user
- 7. Receive an interrupt from the disk I/O subsystem(I/O completed)
- 8. Save the registers and process state for the other user
- 9. Determine that the interrupt was from the disk
- 10. Correct the page table and other tables to

showpageisnow in memory

11. Wait for the CPU to be allocated to this process again
12. Restore the user registers, process state, and newpagetable, and then resume the interrupted instruction



Operating System Concepts – 10th Edition 10.19 Silberschatz, Galvin and Gagne ©2018



Performance of Demand Paging

Three major activities

- Service the interrupt – careful coding means just several hundred instructions needed
- Input the page from disk – lots of time
- Restart the process – again just a small amount of time

- Page Fault Rate $0 \leq p \leq 1$
 - if $p = 0$ no page faults
 - if $p = 1$, every reference is a fault
- Effective Access Time (EAT)

$$\text{EAT} = (1 - p) \times \text{memory access}$$

+ p (page fault overhead

+ swap page out

+ swap page in)



Operating System Concepts – 10th Edition 10.20 Silberschatz, Galvin and Gagne ©2018



Demand Paging Example

■ Memory access time = 200 nanoseconds ■ Average page-fault service time = 8 milliseconds ■

$$\text{EAT} = (1 - p) \times 200 + p (8 \text{ milliseconds})$$

$$= (1 - p) \times 200 + p \times 8,000,000$$

$$= 200 + p \times 7,999,800$$

- If one access out of 1,000 causes a page fault, then EAT = 8.2 microseconds.

This is a slowdown by a factor of 40!!

- If want performance degradation < 10 percent
 - $220 > 200 + 7,999,800 \times p$
 - $20 > 7,999,800 \times p$
 - $p < .0000025$

4 one page fault in every 400,000 memory accesses



Operating System Concepts – 10th Edition 10.21 Silberschatz, Galvin and Gagne ©2018



Copy-on-Write

- **Copy-on-Write** (COW) allows both parent and child processes to initially **share** the same pages in memory
 - If either process modifies a shared page, only then is the page copied
- COW allows more efficient process creation as only modified pages are copied
- In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages

- Pool should always have free frames for fast demand page execution

4 Don't want to have to free a frame as well as other processing on page fault

- Why zero-out a page before allocating it?
- `vfork()` variation on `fork()` system call has parent suspend and child using copy-on-write address space of parent
 - Designed to have child call `exec()`
 - Very efficient



Operating System Concepts – 10th Edition 10.22 Silberschatz, Galvin and Gagne ©2018



Before Process 1 Modifies Page C



Operating System Concepts – 10th Edition 10.23 Silberschatz, Galvin and Gagne ©2018



After Process 1 Modifies Page C



Operating System Concepts – 10th Edition 10.24 Silberschatz, Galvin and Gagne ©2018



What Happens if There is no Free Frame?

- Used up by process pages
- Also in demand from the kernel, I/O buffers, etc
- How much to allocate to each?
- Page replacement – find some page in memory, but not really in use, page it out
 - Algorithm – terminate? swap out? replace the page?
 - Performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times



- Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- Use **modify (dirty) bit** to reduce overhead of page transfers—only modified pages are written to disk
- Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory





Need For PageReplacement





Basic Page Replacement

1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Note now potentially 2 page transfers for page fault – increasing EAT





Page Replacement





Page and Frame Replacement Algorithms

- **Frame-allocation algorithm** determines
 - How many frames to give each process
 - Which frames to replace
- **Page-replacement algorithm**
 - Want lowest page-fault rate on both first access and re-access
 - Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - String is just page numbers, not full addresses
 - Repeated access to the same page does not cause a page fault
 - Results depend on number of frames available
- In all our examples, the **reference string** of referenced page numbers is

7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0
,1,7,0,1



Operating System Concepts – 10th Edition 10.30 Silberschatz, Galvin and Gagne ©2018



Graph of Page Faults Versus the Number of Frames





First-In-First-Out (FIFO) Algorithm.

Reference string: 7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1

- 3 frames (3 pages can be in memory at a time per process)



15 page faults

- How to track ages of pages?
 - Just use a FIFO queue



Belady's Anomaly

Consider the string
1,2,3,4,1,2,5,1,2,3,4,5 • Adding more frames can cause more page faults! •
Graph illustrating Belady

's Anomaly



Optimal Algorithm

- Replace page that will not be used for longest period of time • 9 is optimal for the example



- How do you know this?
 - Can't read the future
 - Used for measuring how well your algorithm performs



- Optimal is an example of **stack algorithms** that don't suffer from

Belady's Anomaly





Least Recently Used(LRU) Algorithm

■ Use past knowledge rather than future ■ Replace page that has not been used in the most amount of time ■ Associate time of last use with each page



- 12 faults – better than FIFO but worse than OPT
- Generally good algorithm and frequently used
- LRU is another example of stack algorithms; thus it doesnot sufferfrom Belady's Anomaly





LRU Algorithm Implementation.

Time-counter implementation

- Every page entry has a time-counter variable; every time a page is referenced through this entry, copy the value of the clock into the time-counter
- When a page needs to be changed, look at the time-counters to find smallest value

4 Search through a table is needed

■ Stack implementation

- Keep a stack of page numbers in a double link form: • Page referenced:

4 Move it to the top

4 Requires 6 pointers to be changed

- But each update more expensive
- No search for replacement





Stack Implementation

Use of a stack to record most recent page references



LRU Approximation Algorithms.

Needs special hardware

▪ **Reference bit**

- With each page associate a bit, initially = 0
- When page is referenced bit set to 1
- Replace any with reference bit = 0 (if one exists)
- We do not know the order, however





LRU Approximation Algorithms(cont.)

Second-chance algorithm • Generally FIFO, plus hardware-provided reference bit ■ **Clock** replacement

- If page to be replaced has
 - 4 Reference bit = 0 -> replace it
 - 4 Reference bit = 1 then:
 - Set reference bit 0, leave page in memory
 - Replace next page, subject to same rules





Second-chanceAlgorithm





Enhanced Second-Chance Algorithm

- Improve algorithm by using reference bit and modify bit (if available) in concert
- Take ordered pair (reference, modify):
 - (0, 0) neither recently used nor modified – best page to replace
 - (0, 1) not recently used but modified – not quite as good, must write out before replacement
 - (1, 0) recently used but clean – probably will be used again soon
 - (1, 1) recently used and modified – probably will be used again soon and need to write out before replacement
- When page replacement called for, use the clock scheme but use the four classes to replace page in lowest non-empty class • Might need to search circular queue several times





Counting Algorithms

- Keep a counter of the number of references that have been made to each page
 - Not common
- **Least Frequently Used (LFU) Algorithm:**
 - Replaces page with smallest count
- **Most Frequently Used (MFU) Algorithm:**
 - Based on the argument that the page with the smallest count was probably just brought in and has yet to be used





Page-Buffering Algorithms.

Keep a pool of free frames, always

- Then frame available when needed, not found at fault time • Read page into free frame and select victim to evict and add to free pool
- When convenient, evict victim ■ Possibly, keep list of modified pages
- When backing store otherwise idle, write pages there and set to non-dirty
- Possibly, keep free frame contents intact and note what is in them • If referenced again before reused, no need to load contents again from disk
 - Generally useful to reduce penalty if wrong victim frame selected





Page-Buffering Algorithms

- Keep a pool of free frames which is never empty
- Thus a frame is available when needed, not found at fault time
- Read page into free frame and select victim to evict and add to free pool
 - When convenient, evict victim
 - Possibly, keep list of modified pages
 - When backing store otherwise idle, write pages there and set to non-dirty
- Possibly, keep free frame contents intact and note what is in them
- If referenced again before reused, no need to load contents again from disk
 - Generally useful to reduce penalty if wrong victim frame selected





Applications and Page Replacement.

- All of these algorithms have OS guessing about future page access
- Some applications have better knowledge – i.e., databases
 - Memory intensive applications can cause double buffering
 - OS keeps copy of page in memory as I/O buffer
 - Application keeps page in memory for its own work
 - Operating system can provide direct access to the disk, getting out of the way of the applications
 - **Raw disk** mode
 - Bypasses buffering, locking, etc.





Allocation of Frames

Each process needs *minimum* number of frames ■ Example: IBM 370 – 6 pages to handle SS

MOVE instruction:

- Instruction is 6 bytes, might span 2 pages
- 2 pages to handle *from*
- 2 pages to handle *to*
- **Maximum** of course is total frames in the system
- Two major allocation schemes
 - Fixed allocation
 - Priority allocation
- Many variations





- ## Fixed Allocation
- Equal allocation – For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames
 - Keep some as free frame buffer pool
 - Proportional allocation – Allocate according to the size of process
 - Dynamic as degree of multiprogramming, process sizes change

$$m = 64$$

$$s_p =$$

size of process

i

$$s_1 = 10$$

$$S_s = \sum$$

i

$$s_2 = 127 \quad a_1 = 10$$

$$m =$$

total number of frames $137 \times 62 \approx 4$

m_{ss}

i

~~$a_p =$~~ \times allocation for

$i i$

$a_2 = 127$

$137 \times 62 \approx 57$



Operating System Concepts – 10th Edition 10.47 Silberschatz, Galvin and Gagne ©2018



Global vs. Local Allocation

- **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another • Process execution time can vary greatly
 - Greater throughput so more commonly used
- **Local replacement** – each

process selects from only its own set of allocated frames

- More consistent per-process performance
- But possibly underutilized memory
- What if a process does not have enough frames?



Operating System Concepts – 10th Edition 10.48 Silberschatz, Galvin and Gagne ©2018



Reclaiming Pages

- A strategy to implement global page-replacement policy
- All memory requests are satisfied from the free-frame list, rather than waiting for the list to drop to zero before we begin selecting pages for replacement,
- Page replacement is triggered when the list falls

below a certain threshold.

- This strategy attempts to ensure there is always sufficient free memory to satisfy new requests.



Operating System Concepts – 10th Edition 10.49 Silberschatz, Galvin and Gagne ©2018



Reclaiming Pages Example





Operating System Concepts – 10th Edition 10.50 Silberschatz, Galvin and Gagne ©2018




Non-Uniform Memory Access ■ So far, we

assumed that all memory accessed equally ■ Many systems are **NUMA** – speed of access to memory varies

- Consider system boards containing CPUs and memory, interconnected over a system bus
- NUMA multiprocessing architecture



Operating System Concepts – 10th Edition 10.51 Silberschatz, Galvin and Gagne ©2018



NUMA Access(Cont.) ■ Optimal performance comes

from allocating memory

“close to” the

CPU on which the thread is scheduled

- And modifying the scheduler to schedule the thread on the same system board when possible
 - Solved by Solaris by creating **lggroups**
- 4 Structure to track CPU / Memory low latency groups 4 Used my schedule and pager
- 4 When possible schedule all threads of a process and allocate all memory for that process within the lgroup





Thrashing

- If a process does not have

“enough” pages, the page-fault rate

is very high

- Page fault to get page
- Replace existing frame
- But quickly need the replaced frame back ■ This leads to:
- Low CPU utilization
- Operating system thinking that it needs to increase the degree of multiprogramming
- Another process added to the system





Thrashing(Cont.) ■ **Thrashing.** A process is busy swapping pages in and out





Demand Paging and Thrashing

■ Why

does demand paging work?

Locality model

- Process migrates from one locality to another
- Localities may overlap
- Why does thrashing occur?


Σ size of locality > total memory size

- To avoid thrashing:
 - Calculate the Σ size of locality
 - Policy:
 - 4 if Σ size of locality > total memory size → suspend or swap out one of the processes
- Issue: how to calculate “ Σ size of locality”





Locality In A Memory-Reference Pattern³⁴




32


30

28

x
x
x
r
x
x



Operating System Concepts – 10th Edition 10.56 Silberschatz, Galvin and Gagne ©2018




Working-Set Model


- $\Delta \equiv$ working-set window \equiv a fixed number of page references Example: 10,000 instructions
- WSS_i (working set of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)
 - if Δ too small will not encompass the entire locality
 - if Δ too large will

encompass several localities • if $\Delta = \infty \Rightarrow$ will encompass entire program

■ Example




Operating System Concepts – 10th Edition 10.57 Silberschatz, Galvin and Gagne ©2018



Working-Set Model

- $\Delta \equiv$ working-set window \equiv a fixed number of page references Example: 10,000 instructions
- WSS_i (working set of Process P_i) = total number of pages referenced in the most recent Δ (varies in time)
 - if Δ too small will not encompass the entire locality • if Δ too large will

encompass several localities • if $\Delta = \infty \Rightarrow$ will encompass entire program



■ $D = \Sigma$

$WSS_i \equiv$ total demand frames • Approximation of locality



Operating System Concepts – 10th Edition 10.58 Silberschatz, Galvin and Gagne ©2018



Working-Set Model (Cont.)

- $D = \Sigma WSS_i \equiv$ total demand frames
 - Approximation of locality
- m = total number of frames
- If $D > m \Rightarrow$ Thrashing
- Policy if $D > m$, then suspend or swap out one of the processes



Keeping Track of the Working Set

- Approximate with interval timer + a reference bit
- Example: $\Delta = 10,000$
 - Timer interrupts after every 5000 time units
 - Keep in memory 2 bits for each page i – $B1_i$ and $B2_i$
 - Whenever a timer interrupts copy the reference to one

of the B_j and sets the values of all reference bits to 0. If either $B1_i$ or $B2_i = 1$, it implies that Page i is in the working set. Why is this not completely accurate?

- Improvement = 10 bits and interrupt every 1000 time units



Operating System Concepts – 10th Edition 10.60 Silberschatz, Galvin and Gagne ©2018



Working Sets and Page Fault Rates

- Direct relationship between working set of a process and its page-fault rate
- Working set changes over time
- Peaks and valleys over time



Operating System Concepts – 10th Edition 10.61 Silberschatz, Galvin and Gagne ©2018



Page-Fault Frequency Algorithm.

More direct approach than WSS

- Establish “acceptable” **page-fault frequency (PFF)** rate and use local replacement policy
 - If actual rate too low, process loses frame
 - If actual rate too high,

process gains frame



Operating System Concepts – 10th Edition 10.62 Silberschatz, Galvin and Gagne ©2018



Allocating Kernel Memory

▪ Treated differently from user memory

- Often allocated from a free-memory pool
 - Kernel requests memory for structures of varying sizes
 - Some kernel

memory needs to be contiguous ⁴ i.e., for device I/O

- Two schemes:
 - Buddy System
 - Slab Allocator



Operating System Concepts – 10th Edition 10.63 Silberschatz, Galvin and Gagne ©2018



Buddy System ■ Allocates memory from fixed-size segment consisting of physically-contiguous pages

- Memory allocated using **power-of-2 allocator** • Satisfies requests in units sized as power of 2 • Request rounded up to next highest power of 2 • When smaller allocation needed than is available, current chunk split into

two buddies of next-lower power of 2

4 Continue until appropriate sized chunk available



Operating System Concepts – 10th Edition 10.64 Silberschatz, Galvin and Gagne ©2018



Buddy System Example

Assume 256KB chunk available, kernel requests 21KB • Split into A_L and A_R of 128KB each

4 One further divided into B_L and B_R of 64KB

- One further into C_L and C_R of 32KB each – one used to satisfy request
- Advantage – quickly **coalesce** unused chunks into larger chunk
- Disadvantage - fragmentation

Operating System Concepts – 10th Edition 10.65 Silberschatz, Galvin and Gagne ©2018



Buddy System Allocator







Slab Allocator

▪ Alternate strategy

- **Slab** is one or more physically contiguous pages
- **Cache** consists of one or more slabs
- Single cache for each unique kernel data structure
- Each cache filled with **objects** – instantiations of the data structure
- When cache created, filled with objects marked as **free**
- When structures stored, objects marked as **used**
- If slab is full of used objects, next object allocated from empty slab
- If no empty slabs, new slab allocated
- Benefits include no fragmentation, fast memory request satisfaction





Slab Allocation





Slab Allocator in Linux

For example, process descriptor is of type `struct task_struct` ■ Approximately 1.7 KB of memory

- New task -> allocate new struct from cache • Will use existing free `struct task_struct`
- Slab can be in three possible states

1. Full – all used
2. Empty – all free
3. Partial – mix of free and used

- Upon request, slab allocator
 1. Uses free struct in partial slab
 2. If none, takes one from empty slab
 3. If no empty slab, create new empty





Slab Allocator inLinux(Cont.)

Slab started in Solaris, now wide-spread for both kernel mode and user memory in various OSes

- Linux 2.2 had SLAB, now has both SLOB and SLUBAllocators
 - SLOB for systems with limited memory
 - 4 Simple List of Blocks – maintains 3 list objects for small, medium, large objects
 - SLUB is performance-optimized SLAB removes per- CPU queues, metadata stored in page structure





Other Considerations.

Prepaging

- Page size
- TLB reach
- Inverted page table
- Program structure
- I/O interlock and page locking





Prepaging

- To reduce the large number of page faults that occurs at process startup
- Prepage all or some of the pages a process will need, before they are referenced
- But if prepaged pages are unused, I/O and memory was wasted
- Assume s pages are prepaged and α of the pages is used • Question: is the cost of $s * \alpha$ save pages faults is greater or less than the cost of prepaging $s * (1 - \alpha)$ unnecessary pages?
 - If α is close to 0 \Rightarrow prepaging loses
 - If α is close to 1 \Rightarrow prepaging wins





Page Size

Sometimes OS designers have a choice • Especially if running on custom-built CPU ■ Page size selection must take into consideration:

- Fragmentation
 - Page table size
 - **Resolution**
 - I/O overhead
 - Number of page faults
 - Locality
 - TLB size and effectiveness
- Always power of 2, usually in the range 2^{12} (4,096 bytes) to 2^{22} (4,194,304 bytes)
 - On average, growing over time





TLB Reach

- TLB Reach - The amount of memory accessible from the TLB
- TLB Reach = (TLB Size) X (Page Size)
- Ideally, the working set of each process is stored in the TLB
- Otherwise there is a high degree of page faults
- Increase the Page Size
 - This may lead to an increase in fragmentation as not all applications require a large page size
- Provide Multiple Page Sizes
 - This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation





ProgramStructure

- Program structure

- `int[128,128] data;`
- Each row is stored in one page
- Program 1

```
for (j = 0; j < 128; j++)  
for (i = 0; i < 128; i++)  
    data[i,j] = 0;
```

128 x 128 = 16,384 page faults

- Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i,j] = 0;
```

128 page faults



I/Ointerlock

- **I/O Interlock** – Pages must sometimes be locked into
 - Consider I/O - Pages that are used for copying a file from a device must be locked from being selected eviction by a page replacement algorithm
- **Pinning** of pages to lock into memory



memory

for

Operating System Concepts – 10th Edition 10.76 Silberschatz, Galvin and Gagne ©2018



Operating System Examples

- Solaris

- Windows



Windows

- Uses demand paging with **clustering**. Clustering brings in pages surrounding the faulting page
- Processes are assigned **working set minimum** and **working set maximum**



- Working set minimum is the minimum number of pages the process is guaranteed to have in memory
- A process may be assigned as many pages up to its working set maximum
- When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- Working set trimming removes pages from processes that have pages in excess of their working set minimum



Operating System Concepts – 10th Edition 10.78 Silberschatz, Galvin and Gagne ©2018



Solaris

- Maintains a list of free pages to assign faulting processes
- **Lotsfree** – threshold parameter (amount of free memory) to begin paging
- **Desfree** – threshold parameter to increasing paging
- **Minfree** – threshold

- parameter to being swapping
- Paging is performed by **pageout** process
- **Pageout** scans pages using modified clock algorithm
- **Scanrate** is the rate at which pages are scanned. This ranges from **slowscan** to **fastscan**
- **Pageout** is called more frequently depending upon the amount of free memory available
 - **Priority paging** gives priority to process code pages



Operating System Concepts – 10th Edition 10.79 Silberschatz, Galvin and Gagne ©2018



Solaris 2 PageScanner

